

## (1) Data Sampling:

### Step 1: Down-sample the majority class

- Omit many of the majority class examples from the training data to make the dataset more balanced

### Step 2: Upweight the down-sampled class

- Down-sampling introduces “predictive bias” by showing the model an artificial world where the classes are more balanced than the real world
- To solve this, we “upweight” the loss in the majority class by the factor by which we have down-sampled the majority class

Idea Reference: <https://developers.google.com/machine-learning/crash-course/overfitting/imbalance-datasets>

## (2) Transactional Graph Construction (using Pytorch-geometric):

- Splitting of the dataset into train-test-validation temporally so that future information is not leaked during the training
- Performed SMOTE-like oversampling by duplicating fraud transactions with slight noise in training dataset keeping validation and test dataset intact to increase the number of fraud transactions in the training dataset

## (3) Graph Transformers Model (i.e. GAT model):

Link to the paper on GAT : <https://arxiv.org/pdf/1710.10903.pdf>

The graph attentional operator from the “Graph Attention Networks” paper.

$$\mathbf{x}'_i = \sum_{j \in \mathcal{N}(i) \cup \{i\}} \alpha_{i,j} \Theta_t \mathbf{x}_j,$$

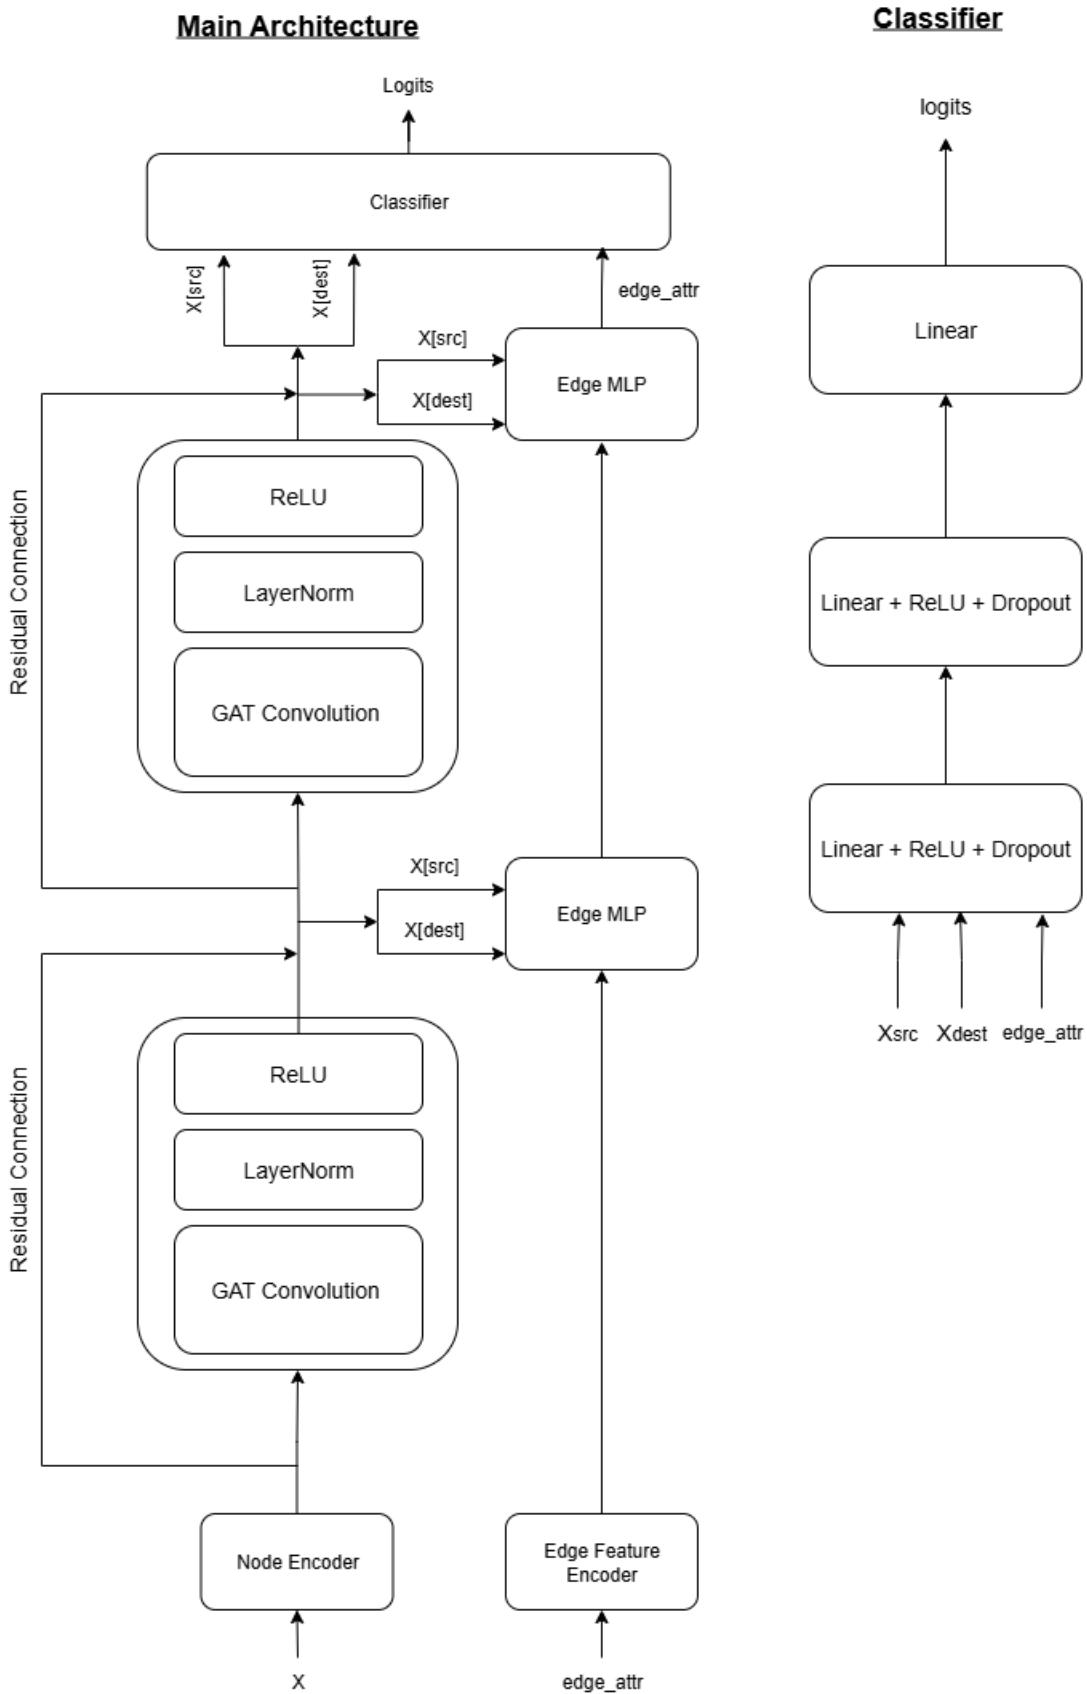
where the attention coefficients  $\alpha_{i,j}$  are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}_s^\top \Theta_s \mathbf{x}_i + \mathbf{a}_t^\top \Theta_t \mathbf{x}_j))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_s^\top \Theta_s \mathbf{x}_i + \mathbf{a}_t^\top \Theta_t \mathbf{x}_k))}.$$

If the graph has multi-dimensional edge features  $\mathbf{e}_{i,j}$ , the attention coefficients  $\alpha_{i,j}$  are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}_s^\top \Theta_s \mathbf{x}_i + \mathbf{a}_t^\top \Theta_t \mathbf{x}_j + \mathbf{a}_e^\top \Theta_e \mathbf{e}_{i,j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_s^\top \Theta_s \mathbf{x}_i + \mathbf{a}_t^\top \Theta_t \mathbf{x}_k + \mathbf{a}_e^\top \Theta_e \mathbf{e}_{i,k}))}.$$

If the graph is not bipartite,  $\Theta_s = \Theta_t$ .



## Training:

$$(1) \text{ BCE Loss} = -[pos_{weight} * y * \log(p) + (1 - y) * \log(1 - p)]$$

Where:

$y$  = true label(0 or 1)

$p$  = predicted probability

$pos_{weight} = class\_count[0] / class\_count[1]$

$$(2) \text{ Focal Loss} = -[\alpha * (1 - p)^{\gamma} * \log(p)]$$

Where:

$\alpha$  (alpha): Balances positive/negative examples ( $= 1 - (fraud\_count) / total$ )

$\gamma$  (gamma): Focuses learning on hard examples ( $= 2$ )

$p$ : Model's confidence for the correct class

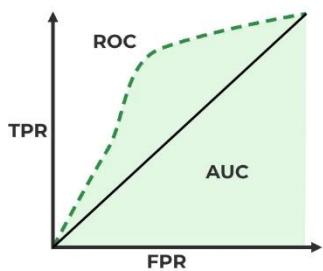
$$(3) \text{ BCE with logits} = \begin{cases} \log(1 + \exp(-\text{logit})) = \log(\text{sigmoid}(-\text{logit})), & \text{target} = 1 \\ \log(1 + \exp(+\text{logit})) = \log(\text{sigmoid}(\text{logit})), & \text{target} = 0 \end{cases}$$

$$(4) \text{ Total Loss} = 0.7 * \text{Focal Loss} + 0.3 * \text{BCE Loss}$$

## Evaluation:

$$(1) \text{ ROC-AUC curve (Receiver Operating Characteristic Area under the curve):}$$

- Curve of TPR(true positive rate) vs FPR(false positive rate)
- $TPR = TP / (TP + FN)$ ,  $FPR = FP / (FP + TN)$
- Ideally, TPR should be 1 and FPR should be 0



$$(2) \text{ Precision} = TP / (TP + FP)$$

$$(3) \text{ Recall} = TP / (TP + FN)$$

$$(4) \text{ F1-Score} = (2 * \text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$